

```
#Problem 3: Regression Estimator  
# For this problem, the dataset to be used is the diabetes dataset  
from sklearn.
```

```
import pandas as pd  
import numpy as np  
from sklearn.datasets import load_diabetes  
from sklearn.model_selection import train_test_split, GridSearchCV  
from sklearn.linear_model import LinearRegression, Ridge, Lasso  
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Load the diabetes dataset
```

```
diabetes = load_diabetes()  
X = diabetes.data  
y = diabetes.target
```

```
# Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
# Define the regressors to be tested
```

```
regressors = {'Linear Regression': LinearRegression(),  
              'Ridge': Ridge(),  
              'Lasso': Lasso()}
```

```
# Define the parameter grid for the Ridge regressor
```

```
param_grid = {'alpha': [0.1, 1, 10]}
```

```
# Train and evaluate the regressors
```

```
results = {}
```

```
for name, regressor in regressors.items():
```

```
    print('Training', name)
```

```
    # For Ridge, perform a grid search to select the best alpha  
parameter
```

```
    if name == 'Ridge':
```

```
        grid_search = GridSearchCV(regressor, param_grid, cv=5)
```

```
        grid_search.fit(X_train, y_train)
```

```
        best_regressor = grid_search.best_estimator_
```

```
    else:
```

```
        best_regressor = regressor.fit(X_train, y_train)
```

```
# Make predictions on the testing set
```

```
y_pred = best_regressor.predict(X_test)
```

```
# Calculate performance metrics
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
rmse = np.sqrt(mse)
```

```
r2 = r2_score(y_test, y_pred)
```

```

# Save the results
results[name] = {'Regressor': best_regressor,
                 'MSE': mse,
                 'RMSE': rmse,
                 'R2': r2}

# Print the results
for name, result in results.items():
    print(name, 'performance:')
    print('MSE:', result['MSE'])
    print('RMSE:', result['RMSE'])
    print('R2:', result['R2'])
    print('Regressor:', result['Regressor'])
    print()

#Plot the results
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

# Plot the results
fig, ax = plt.subplots(1, 3, figsize=(15, 5))
for i, name in enumerate(results.keys()):
    ax[i].scatter(y_test, results[name]['Regressor'].predict(X_test))
    ax[i].set_title(name)
    ax[i].set_xlabel('True')
    ax[i].set_ylabel('Predicted')
    # Plot the best fit line
    x = np.linspace(0, 300, 100)
    ax[i].plot(x, x, color='red')

plt.show()

# Select the best regressor based on RMSE
best_regressor = None
best_rmse = float('inf')
for name, result in results.items():
    if result['RMSE'] < best_rmse:
        best_regressor = result['Regressor']
        best_rmse = result['RMSE']
print('Best regressor based on RMSE:', best_regressor)

```

```

Training Linear Regression
Training Ridge
Training Lasso
Linear Regression performance:
MSE: 2900.19362849348
RMSE: 53.853445836765914
R2: 0.4526027629719197

```

Regressor: LinearRegression()

Ridge performance:

MSE: 2856.486887670654

RMSE: 53.446111997699646

R2: 0.46085219464119254

Regressor: Ridge(alpha=0.1)


Lasso performance:

MSE: 3403.5757216070733

RMSE: 58.340172450954185

R2: 0.3575918767219115

Regressor: Lasso()



Best regressor based on RMSE: Ridge(alpha=0.1)